

- Quantizer
 - n-level -> split into n intervals
- Uniform
 - $Q(x) = \lfloor \frac{x-d_0}{\Delta} \rfloor$ // this gives i where $x \in [d_i, d_{i+1})$ //Exercise: Prove that
 - $Q^{-1}(i) = d_0 + \left(i + \frac{1}{2}\right)\Delta$ // that is because $Q^{-1}(i) \stackrel{\text{def}}{=} r_i$ and $r_i = d_0 + \left(i + \frac{1}{2}\right)\Delta$
- Non-uniform
 - q: binary search
 - deq: use r_i
- Semi uniform
 - **Quantization:** as in uniform quantizers, $Q(x) = \lfloor \frac{x-d_0}{\Delta} \rfloor$, which takes $O(1)$ time
 - **Dequantization:** $Q^{-1}(i) = r_i$, which takes a constant time ($O(1)$)
- Max LLOYD
 - NON-UNIFORM
 - Initial the value as uniform q
 - loop
 - r_i = average in $[d_i, d_{i+1})$
 - $d_i = (r_{i-1} + r_i)/2$
 - until d_i doesn't change much
- Transform
 - **Specifically, the transforms must have the following properties:**
 - **Decorrelation of data**
 - **Separation of data** into vision-sensitive data and vision-insensitive data
 - **Energy compaction:** concentrating the important data into a very small subset
 - **Invertability:** since data loss should occur only in quantization, transforms must be lossless
 - decorrelated data: less blurring, less loss of patterns.

• Example:
$$\begin{bmatrix} 1 & 5 \\ 3 & 4 \\ 2 & 1 \end{bmatrix} \cdot \begin{bmatrix} 2 & 4 & 1 & 3 \\ 1 & 2 & 4 & 5 \end{bmatrix} = \begin{bmatrix} 7 & 14 & 21 & 28 \\ 10 & 20 & 19 & 29 \\ 5 & 10 & 6 & 11 \end{bmatrix}$$

THE MATRIX OF THE HADAMARD TRANSFORM

- To express a_{kl} , we need to express k and l in binary (using $n = \log N$ bits)

$$k = k_{n-1}k_{n-2} \dots k_1k_0, \quad l = l_{n-1}l_{n-2} \dots l_1l_0$$

$$a_{kl} = \frac{1}{\sqrt{N}} (-1)^{k_{n-1}l_{n-1} + k_{n-2}l_{n-2} + \dots + k_1l_1 + k_0l_0}$$

$$A_N^{-1} = A_N^T = A_N$$

- Illustrations:

$$A_2 = \sqrt{\frac{1}{2}} \begin{pmatrix} 1 & 1 \\ 1 & -1 \end{pmatrix} \quad A_4 = \frac{1}{2} \begin{pmatrix} 1 & 1 & 1 & 1 \\ 1 & -1 & 1 & -1 \\ 1 & 1 & -1 & -1 \\ 1 & -1 & -1 & 1 \end{pmatrix}$$

$$A_8 = \sqrt{\frac{1}{8}} \begin{pmatrix} 1 & 1 & 1 & 1 & 1 & 1 & 1 & 1 \\ 1 & -1 & 1 & -1 & 1 & -1 & 1 & -1 \\ 1 & 1 & -1 & -1 & 1 & 1 & -1 & -1 \\ 1 & -1 & -1 & 1 & 1 & -1 & -1 & 1 \\ 1 & 1 & 1 & 1 & -1 & -1 & -1 & -1 \\ 1 & -1 & 1 & -1 & -1 & 1 & -1 & 1 \\ 1 & 1 & -1 & -1 & -1 & -1 & 1 & 1 \\ 1 & -1 & -1 & 1 & -1 & 1 & 1 & -1 \end{pmatrix}$$

CS6381 Data Compression

Transforms- Part I

AN ALTERNATIVE DEFINITION OF THE MATRIX OF THE HADAMARD TRANSFORM

- The Hadamard matrix can be defined recursively (where N is a power of 2, i.e., $N = 2^n$ for some positive integer n)

$$A_1 = (1) \text{ and } \forall N > 1, A_N = \sqrt{\frac{1}{2}} \begin{pmatrix} A_{\frac{N}{2}} & A_{\frac{N}{2}} \\ A_{\frac{N}{2}} & -A_{\frac{N}{2}} \end{pmatrix}$$

- **Exercise:** Derive A_2, A_4 and A_8 using this recursive definition
- **Exercise:** Compare the values A_2, A_4 and A_8 with their values on the previous slide to verify that the two definitions are equivalent
- **Exercise:** Using the recursive definition of A_N , prove by induction on n that A_N is a symmetric matrix, and that $A_N A_N = I_N$, i.e., $A_N^{-1} = A_N$.
- **Note:** In Matlab, if you call `hadamard(N)`, it returns to you the Hadamard matrix A_N but without the constant multiplier $\sqrt{\frac{1}{N}}$

SPEED OF THE TRANSFORM

-- 2D SIGNALS --

- For 2D input signals, i.e., an $N \times M$ image X
- Recall that $Y = A_N X A_M^T$ (i.e., we transform every column then every row)
- Every column of A takes $O(N^2)$ time to compute, and so the M columns take $O(MN^2)$ time
- Every row takes $O(M^2)$ time to compute, and so the N columns take $O(NM^2)$ time
- Thus, the 2D transform takes $O(MN^2 + NM^2) = O(N^3)$ for $N = M$
- For $N=1000$, $O(N^3)=1$ second on a 1GFLOPS device (or 17 mins on a 1MFLOPS)
- But even on a 1GFLOPS device, applying the 2D transform on all the frames of a 2-hour video (with a 30fps rate) takes 2.5 days!

29

- Lossless
 - DPCM

DPCM

-- MAIN METHOD (FOR 1D DATA): PARAMETER --

- The choice of the parameter a is up to the user, and can be adapted to your data/apps
- It can also be optimized to fit your input data:
 - Let $E(a) = \sum_{i=1}^n e_i^2 = \sum_{i=1}^n (x_i - ax_{i-1})^2$, assuming that $x_0 = 0$
 - The smaller the magnitudes of the residuals e_i , the better, since that would increase the redundancy (and make certain residual values occur with high probabilities)
 - Therefore, we are interested in minimizing $E(a)$
 - To get the optimal value of a that minimizes $E(a)$, compute the derivative $E'(a)$, set it to 0, and solve for a
 - $E'(a) = -2 \sum_{i=1}^n x_{i-1} (x_i - ax_{i-1}) = -2 [\sum_{i=1}^n x_{i-1} x_i - a \sum_{i=1}^n x_{i-1}^2]$,

- $\text{NINT}(z)$ = nearest integer to z
 - Ex: $\text{NINT}[1.7]=2$,
 - $\text{NINT}[1.3]=1$
 - $\text{NINT}[1.5]=2$
 - Setting $E'(a) = 0$, and solving, we get $a = \frac{\sum_{i=1}^n x_{i-1} x_i}{\sum_{i=1}^n x_{i-1}^2}$
 - If x is integers and you want to keep the residues integer, take $a = \text{NINT}(\frac{\sum_{i=1}^n x_{i-1} x_i}{\sum_{i=1}^n x_{i-1}^2})$

30

CS6381 Data Compression

Lossless Compression Part III

DPCM

-- MAIN METHOD: FOR 2D DATA --

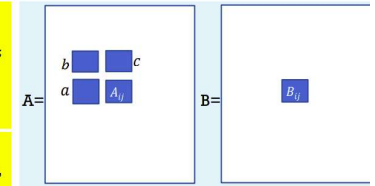
- DPCM for 2D data (i.e., an image A)
- DPMC for 2D data takes three parameters (a, b, c)

Coder Method (to code image A):

1. Compute a residual image B :
 $B[i,j] = A[i,j] - (aA[i,j-1] + bA[i-1,j-1] + cA[i-1,j]);$
 // any pixel out of boundary is assumed = 0
2. Code the residual B with Huffman (or any other suitable lossless coder like Bitplane Coding to be covered later)

Decoder Method:

1. Decode the coded bitstream with the right decoder, to get the residual image B
2. For $i = 1; n$ do // image is an $n \times m$ matrix.
 For $j = 1; m$ do // pixels out of boundary are = 0
 $A[i,j] = B[i,j] + (aA[i,j-1] + bA[i-1,j-1] + cA[i-1,j]);$



31

CS6381 Data Compression

Lossless Compression Part III

DPCM

-- MAIN METHOD (FOR 2D DATA): PARAMETERS --

- The choice of the parameters (a, b, c) is up to the user, and can be adapted to your data/applications
- It can also be optimized to fit your input data:
 - Let $E(a, b, c) = \sum_{i=1}^n \sum_{j=1}^m B_{ij}^2 = \sum_{i=1}^n \sum_{j=1}^m [A[i,j] - (aA[i,j-1] + bA[i-1,j-1] + cA[i-1,j])]^2$
 - Compute the partial derivatives of E with respect to a, b , and c , set them to 0, and solve the system of 3 linear equations with three unknowns (the unknowns are a, b , and c)
- **Exercise:** Carry out the above step, and derive formulas for optimal a, b , and c .
- Note: If you would like to keep every thing integer (like the pixels in image B), take the NINT values of the optimal values of a, b , and c .

- Gray codes

- make adjacent binary only change 1 number

GRAY CODES

-- CONSTRUCTION EXAMPLES --

- Let $G_n = 0G_{n-1}, 1G_{n-1}^R$, that is concatenate the two sequence $0G_{n-1}$ and $1G_{n-1}^R$
- Construct G_2 from G_1 :
 1. We saw that G_1 is: 0, 1 and so $0G_1$ is 00, 01
 2. G_1^R is: 1, 0 and so $1G_1^R$ is: 11, 10
 3. Hence, $G_2 = 0G_1, 1G_1^R$ is: 00, 01, 11, 10
- Construct G_3 from G_2 :
 1. $G_2: 00, 01, 11, 10 \Rightarrow 0G_2: 000, 001, 011, 010$
 2. $G_2^R: 10, 11, 01, 00 \Rightarrow 1G_2^R: 110, 111, 101, 100$
 3. $G_3 = 0G_2, 1G_2^R$ is: 000, 001, 011, 010, 110, 111, 101, 100
- Construct $G_4 = 0G_3, 1G_3^R$: 0000, 0001, 0011, 0010, 0101, 0111, 0101, 0100, 1100, 1101, 1111, 1110, 1010, 1011, 1001, 1000

32

CS6381 Data Compression

Lossless Compression Part III

GRAY CODES

-- CODING OF INTEGERS--

- Once we have a Gray code G_n , we can code any integer k (i.e, a pixel value between 0 and $2^n - 1$) into the k^{th} binary string in G_n

- Example:

- k : 0, 1, 2, 3, 4, 5, 6, 7
- G_3 : 000, 001, 011, 010, 110, 111, 101, 100

Gray_code(0)=000, Gray_code(2)=011
Gray_code(4)=110, Gray_code(7)=100

- For speed reason, it is convenient to know the Gray code $g_{n-1} \dots g_1 g_0$ of an integer k , directly from k , without having to have all G_n (of 2^n strings)

- Method:** input is k , output is the Gray code of k , denoting it $g_{n-1} \dots g_1 g_0$

- Convert k to regular binary $b_{n-1} \dots b_1 b_0$, using decimal-to-binary conversion
- Set $g_{n-1} = b_{n-1}$
- For $i=0$ to $n-2$ do: $g_i = b_i \text{ XOR } b_{i+1}$;

Recall that $(0 \text{ XOR } 0) = (1 \text{ XOR } 1) = 0$,
and $(1 \text{ XOR } 0) = (0 \text{ XOR } 1) = 1$

- Gray-to-binary conversion (for decoding): to get $b_{n-1} \dots b_1 b_0$ from $g_{n-1} \dots g_1 g_0$

- $b_{n-1} = g_{n-1}$;
- For $i=n-2$ down to 0 do: $b_i = g_i \text{ XOR } b_{i+1}$;

GRAY CODES

-- PROOF OF CORRECTNESS--

Theorem: G_n is a Gray code for all $n \geq 1$, and the last string and 1st string (of G_n) differ by only one bit.

Proof: We will prove by induction on n that every two successive strings in G_n differ by exactly one bit, and so do the 1st and last string.

- Basis step:** $n = 1$. Prove that the theorem is true for $n=1$. Well, G_1 : 0, 1, and so its only two successive strings 0 and 1 differ by exactly one bit.

- Induction step:**

- Assume the theorem is true for $n-1$, i.e., every two successive strings in G_{n-1} differ by exactly one bit, and so do the 1st and last string in G_{n-1} . This is called the *induction hypothesis* (IH)
- Prove that every two successive strings in G_n differ by exactly one bit, and so do the 1st and last strings.
- Take two successive strings W_1 and W_2 in $G_n = 0G_{n-1}, 1G_{n-1}^R$. We have three cases:
 - W_1 and W_2 are in $0G_{n-1}$: so $W_1 = 0V_1$ and $W_2 = 0V_2$ where V_1 and V_2 are two successive strings in G_{n-1} , and so, by the IH, they differ by only one bit. Therefore, since W_1 and W_2 agree in the left most bit (0), they differ in only one bit.
 - W_1 and W_2 are in $1G_{n-1}^R$: the proof is very similar to part (a)

21

- W_1 is the last string in $0G_{n-1}$, and W_2 is the 1st string in $1G_{n-1}^R$. So, $W_1 = 0V_1$ and $W_2 = 1V_2$ but also observe that $V_1 = V_2$ because the first string in G_{n-1}^R is the last string of G_{n-1} . Hence, $W_1 = 0V_1$ and $W_2 = 1V_1$, and so they differ by only the leftmost bit.

- One thing remains: W_1 is the first string of $0G_{n-1}$, and W_2 is the last string in $1G_{n-1}^R$. That is similar to case (c) above. Q.E.D.

99

ARITHMETIC CODING (AC)

-- METHOD --

Input: a binary stream $x = x_1 x_2 \dots x_n$, and a probabilistic model of x

Output: a coded bitstream

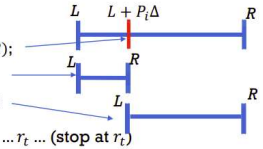
Method:

- Let $I = [L, R)$ where initially $L = 0, R = 1$;
- For $i=1$ to n do

- Let $P_i = \text{Pr}[0/x_1 x_2 \dots x_{i-1}]$; // regardless of what x_i is
- Let $\Delta = R - L$; // length of the interval now
- Split** interval I into 2 subintervals: $[L, L + P_i \Delta)$ and $[L + P_i \Delta, R)$;
- Choose:** If $(x_i = 0)$, reduce I to $[L, L + P_i \Delta)$, i.e., $R := L + P_i \Delta$;
If $(x_i = 1)$, reduce I to $[L + P_i \Delta, R)$, i.e., $L := L + P_i \Delta$;

There is a patent on AC by IBM called the Q-Coder

Observation 1: Every subinterval is nested inside the earlier intervals



- Let $t = \lceil -\log \Delta \rceil$, and $r = \frac{L+R}{2}$ expressed in binary as $0.r_1 r_2 \dots r_t \dots$ (stop at r_t)
- Output: $r_1 r_2 \dots r_t$, i.e., code the input $x_1 x_2 \dots x_n$ as $r_1 r_2 \dots r_t$

Observation 2: r is in every subinterval that was chosen. Why?

$t = \lceil -\log(\text{length of last subinterval}) \rceil$

CS6351 Data Compression

Lossless Compression Part II

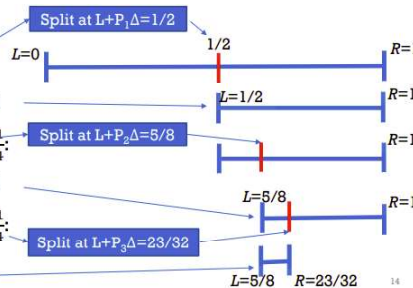
ARITHMETIC CODING (AC)

-- EXAMPLE (1/2) --

- Binary Markov Source $\Pr[0/0] = \Pr[1/1] = \frac{3}{4}$, $\Pr[0/1] = \Pr[1/0] = \frac{1}{4}$, and $\Pr[0] = \Pr[1] = \frac{1}{2}$

- Code input $x = 110$

- $i = 1, x_1 = 1, \Delta = 1, P_1 = \Pr[0] = \frac{1}{2}$
- Since $x_1 = 1$, choose right interval
- $i = 2, x_2 = 1, \Delta = \frac{1}{2}, P_2 = \Pr[0/1] = \frac{1}{4}$
- Since $x_2 = 1$, choose right interval
- $i = 3, x_3 = 0, \Delta = \frac{3}{8}, P_3 = \Pr[0/1] = \frac{1}{4}$
- Since $x_3 = 0$, choose left interval



CS6381 Data Compression

Lossless Compression Part II

ARITHMETIC CODING (AC)

-- EXAMPLE (2/2) --

- The final interval was: $L = \frac{5}{8}, R = \frac{23}{32}$
- Therefore, $L = \frac{5}{8}, R = \frac{23}{32}, \Delta = R - L = \frac{3}{32}$
- $t = \lceil -\log \Delta \rceil = \lceil -\log \frac{3}{32} \rceil = \lceil 3.4150 \rceil = 4$
- So the coded bitstream will have 4 bits
- $\frac{L+R}{2} = \frac{43}{64}$, convert to binary, we get $\frac{L+R}{2} = 0.101011$
- Take the first 4 bits after the point: 1010
- Therefore, the coded bitstream r is: **1010**
- Done with the example

- This example led to expansion, not compression.
- Not surprising because the input is small
- For realistic input, there will be compression, usually $CR \approx 2$

CS6381 Data Compression

Lossless Compression Part II

- decode-> stop at $\text{ceil}(-\log \Delta) = t$
- LZ compression

LEMPER-ZIV (LZ) COMPRESSION

--CODER METHOD--

Input: a binary stream $x = x_1x_2 \dots x_n$ // no need for a probabilistic model

Output: a coded bitstream

Method:

- Maintain a dictionary (DICT) of patterns seen so far
- Set $i=1, j_1=\text{empty}, W=\text{empty}, a=x_1, \text{DICT}[1]=W a$ (i.e., x_1), output $= x_1$;
- Set $i=2$;
- While (there are still input symbols) do
 - Read from the remaining input until the substring scanned (call it S_i) is no longer in DICT
 - Write S_i as $S_i = W_i a_i$, where W_i is in DICT and a_i is the last symbol in S_i (the symbol after W_i)
 - Let j be the index where $W_i = \text{DICT}[j]$; // $j < i$
 - Let $j_1 = \text{d2b}(j)$ using $\lceil \log i \rceil$ bits // d2b is decimal to binary converter
 - Code $W_i a_i$ as (j_1, a_i) and append that code to the output
 - Store $W_i a_i$ in DICT[i], i.e., $\text{DICT}[i] = W_i a_i$;
 - $i++$;

- Observe that:
- $x = \text{DICT}[1] \text{DICT}[2] \text{DICT}[3] \dots$
 - Output = sequence of (j_1, a_i) 's generated

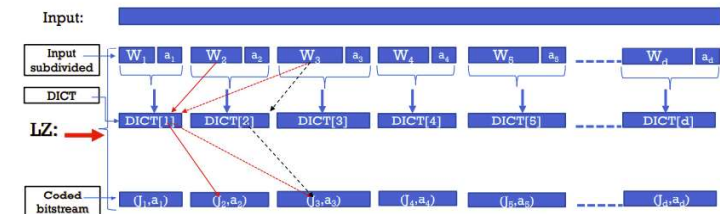
a_i could be non-existent in the very last $W_i a_i$; Code $W_i a_i$ as $(j_1, -)$, i.e., as j_1 alone

CS6381 Data Compression

Lossless Compression Part II

LEMPER-ZIV (LZ) COMPRESSION

--OVERVIEW--



Notes:

- Every W_i is in $\text{DICT}[1 : (i-1)]$
- j_1 is the index where W_i is stored in $\text{DICT}[1 : (i-1)]$, i.e., $W_i = \text{DICT}[j_1]$
- $W_i a_i$ is stored in location i of DICT, i.e., $\text{DICT}[i] = W_i a_i$
- W_i is non-existent, i.e., empty (we denote it as - or ϵ)
- So j_1 is non-existent, i.e., empty (we denote it as -)

CS6381 Data Compression

Lossless Compression Part II

LEMPER-ZIV (LZ) COMPRESSION

-- REMARKS --

- The dictionary is not stored/transmitted in the coded bitstream
- Rather, it is used by the coder to carry out the coding, and then discarded
- As will be seen, the decoder will be able to reconstruct the dictionary from the coded bitstream alone
- No probabilities were needed or used
- The LZ bitrate is asymptotically optimal (i.e., approaches the entropy of the source with memory) **without** the need to know or compute the underlying probability model of the input data.
- The proof will not be provided. See [paper](#) if interested

26

LEMPER-ZIV (LZ) COMPRESSION

-- EXAMPLE --

Input: $x = 0010100100010010100110101$

i	$\lceil \log i \rceil$	$J_i = (j)_2$	W_i	a_i	DICT[i]
1	0	empty	empty	0	0
2	1	$(1)_2 = 1$	0	1	01
3	2	$(10)_2 = 2$	01	0	010
4	2	$(11)_2 = 3$	010	0	0100
5	3	$(100)_2 = 4$	0100	1	01001
6	3	$(101)_2 = 5$	01001	1	010011
7	3	$(011)_2 = 3$	010	1	0101

Notation: When we find $S_i = W_i a_i$ in x , we color W_i in green and a_i in red.
DICT[i] = $W_i a_i$

Output: Sequence of (j, a_i) 's for $i=1,2,\dots$. Therefore:
Output: $(-,0) (1,1) (10,0) (11,0) (100,1) (101,1) (011,1)$. Finally:
Output: 011100110100110110111

Recall:
DICT[0] = empty = $\epsilon = \epsilon$

26

LZ DECODER

-- METHOD --

Input: a coded bitstream $y = y_1 y_2 \dots$; **Output:** the original data x

Method:

- $i=1, W_i = \text{empty}, a_i = y_1, \text{DICT}[1] = W_i a_i$ (i.e., y_1), $x = W_i a_i$ (i.e., y_1);
- $i=2$;
- While (the coded bitstream is not fully scanned) do
 - $j = \text{decimal}(\text{the next } \lceil \log i \rceil \text{ bits from } y)$;
 - $W_i = \text{DICT}[j]$;
 - $a_i = \text{the next symbol from } y$;
 - append $W_i a_i$ to the right of x ;
 - $\text{DICT}[i] = W_i a_i$
 - $i++$;
- // if $a_i = \text{empty}$, at end of decoding, simply append W_i to the right of x

Observe that:

- $x = \text{DICT}[1] \text{ DICT}[2] \text{ DICT}[3] \dots$
- Which is what we want based on what we observed in the LZ coder

27

LZ DECODER

-- EXAMPLE --

Decoding $y = 01110 0110100110110111$ from the previous example

i	$\lceil \log i \rceil$	$j = \text{next } \lceil \log i \rceil \text{ bits of } y$	$W_i = \text{DICT}[j]$	a_i	DICT[i] = $W_i a_i$	$x = (\text{previous}(x))(W_i a_i)$
1	0	empty	empty	0	0	0
2	1	$(1)_2 = 1$	0	1	01	001
3	2	$(10)_2 = 2$	01	0	010	001010
4	2	$(11)_2 = 3$	010	0	0100	0010100100
5	3	$(100)_2 = 4$	0100	1	01001	001010010001001
6	3	$(101)_2 = 5$	01001	1	010011	00101001000100110011
7	3	$(011)_2 = 3$	010	1	0101	00101001000100110101

Decoded data $x = 00101001000100110101$

28

COMPRESSION PERFORMANCE METRICS

-- SIZE REDUCTION MEASURES --

- **Bit Rate** (or **bitrate** or **BR**): Average number of bits per original data element, after compression
 - In text compression: $BR = \frac{\text{avg. number of bits per character}}{\text{characters in the original text}} = \frac{\text{\#bits in the coded bitstream}}{\text{\#characters in the original text}}$
 - In images/videos: $BR = \frac{\text{average number of bits per pixel}}{\text{pixels in the image or video}} = \frac{\text{\#bits in the coded bitstream}}{\text{\#pixels in the image or video}}$
 - In audio: $BR = \frac{\text{avg. num. of bits per sound sample}}{\text{samples in the audio file}} = \frac{\text{\#bits in the coded bitstream}}{\text{\#samples in the audio file}}$
- **Compression Ratio (CR)** = $\frac{\text{size of the uncompressed data (in bits)}}{\text{size of the coded bitstream (in bits)}} = \frac{|I|}{|b|}$
- **Exercise:** A file has 1000 characters, and takes 8K bits before compression. After compression, its size became 4K bits. What is the bitrate? What the CR?

BASIC DEFINITIONS IN COMPRESSION/CODING

-- QUALITY MEASURES: SNR --

- Let I be an original signal (e.g., an image), and \hat{I} be its lossily reconstructed counterpart
- **Signal-to-Noise Ratio (SNR)** in the case of lossy compression:
 - It is a measure of the quality of the reconstructed (decompressed/decoded) data
 - More accurately, it is the fidelity of the decoded data w.r.t. the original
 - Mathematically: $SNR = 10 \log_{10} \left(\frac{\|I\|^2}{\|I - \hat{I}\|^2} \right) = 20 \log_{10} \left(\frac{\|I\|}{\|I - \hat{I}\|} \right)$
 where for any vector/matrix/set of number $E = \{x_1, x_2, \dots, x_N\}$,
 $\|E\|^2 = x_1^2 + x_2^2 + \dots + x_N^2$
- The unit of SNR is "**decibel**" (or **dB** for short)
- So, if $SNR = 23$, we say the SNR is 23 dB

Question: Does it make sense to compute SNR for lossless compression? Why or why not?

BASIC DEFINITIONS IN COMPRESSION/CODING

-- QUALITY MEASURES: MEAN-SQUARE ERROR --

- Mean-Square Error (MSE): $MSE = \frac{1}{N} \|I - \hat{I}\|^2$
- Relative Mean-Square Error (RMSE): $RMSE = \frac{\|I - \hat{I}\|^2}{\|I\|^2}$
- **Therefore**, $SNR = -10 \log_{10} RMSE$
- **Observations:**
 - The smaller the RMSE (or the MSE), the higher the SNR
 - Therefore, the higher the SNR, the better the quality of the reconstructed data
- **Exercise:** Prove that if RMSE is decreased by a factor of 10, then SNR increases by 10 decibels.
- That justifies the multiplicative factor in the definition of the SNR.

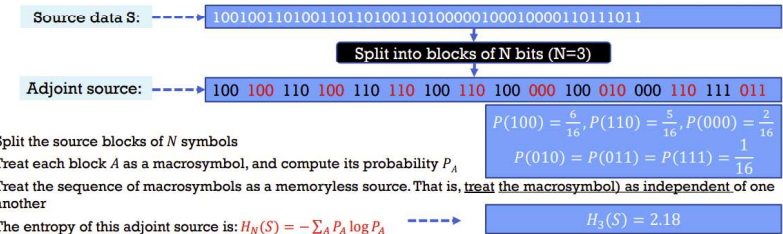
13

- Therefore, **Entropy of S** , denoted $H(S)$, is:

$$H(S) = -(p_1 \log p_1 + p_2 \log p_2 + \dots + p_n \log p_n) = -\sum_{i=1}^n p_i \log p_i$$

INFORMATION THEORY (ADJOINT SOURCES)

- Adjoint Source of Order N of a source with memory:



- Split the source blocks of N symbols
- Treat each block A as a macrosymbol, and compute its probability P_A
- Treat the sequence of macrosymbols as a memoryless source. That is, treat the macrosymbol as independent of one another
- The entropy of this adjoint source is: $H_N(S) = -\sum_A P_A \log P_A$
- **Theorem** (Shannon): $\frac{H_N(S)}{N} \rightarrow H(S)$ as $N \rightarrow \infty$
- This implies that for any source S with memory, if we divide it into blocks of large enough size N and then block-code it without taking advantage of inter-block correlation, then we still approximate the performance of the best coder S .
- **Therefore, block coding is a good way to go when we have correlation**

CS6381 Data Compression

Introduction and Background

43

HUFFMAN CODING

-- THE CODING ALGORITHM --

Input: alphabet $\{a_1, a_2, \dots, a_n\}$ and symbol probabilities $\{p_1, p_2, \dots, p_n\}$

Output: the codewords of the alphabet symbols

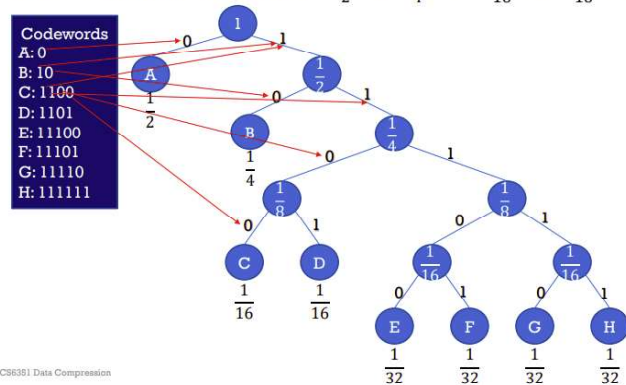
Method: (a Greedy method for creating a Huffman tree as follows)

1. Create a node for each symbol a_i // these nodes will be the leaves
2. **While** (there are two or more uncombined nodes) **do**
 - Select 2 uncombined nodes a and b of minimum probabilities
 - Create a new node c of prob $P_a + P_b$, and make a and b children of c
3. Label the tree edges: left edges with 0, right edges with 1
4. The codeword of each alphabet symbol a_i (a leaf) is the binary string that labels the path from the root to leaf a_i

HUFFMAN CODING

-- ILLUSTRATION OF THE CODING ALGORITHM --

Alphabet = {A, B, C, D, E, F, G, H}, $P_A = \frac{1}{2}$, $P_B = \frac{1}{4}$, $P_C = \frac{1}{16}$, $P_D = \frac{1}{16}$, $P_E = P_F = P_G = P_H = \frac{1}{32}$



HUFFMAN CODING

-- CODING PERFORMANCE --

- In lossless compression of memoryless sources

- If the coder works by computing a codeword for each alphabet symbol
- Then, we can compute a **coder bitrate**, independent of any actual input data

- Notation:

- For any binary string s , denote by $|s|$ the number of bits in s
- Let **codeword**(a_i) denote the codeword for symbol a_i

- **Coder bitrate:** $BR = \sum_{i=1}^n p_i |\text{codeword}(a_i)|$

- **Source Entropy:** $H = -\sum_{i=1}^n p_i \log p_i$

Example: the Huffman coder just presented

- The codewords and the probabilities are

Codewords	Length	Probabilities
A: 0	1	1/2
B: 10	2	1/4
C: 1100	4	1/16
D: 1101	4	1/16
E: 11100	5	1/32
F: 11101	5	1/32
G: 11110	5	1/32
H: 11111	5	1/32

$$BR = 1 \times \frac{1}{2} + 2 \times \frac{1}{4} + 4 \times \frac{1}{16} + 4 \times \frac{1}{16} + 5 \times \frac{1}{32} + 5 \times \frac{1}{32} + 5 \times \frac{1}{32} + 5 \times \frac{1}{32} = \frac{69}{32} = 2.125 \text{ bits/symbol}$$

$$\text{Entropy: } H = -\left(\frac{1}{2} \log \frac{1}{2} + \frac{1}{4} \log \frac{1}{4} + \frac{1}{16} \log \frac{1}{16} + \frac{1}{16} \log \frac{1}{16} + \frac{1}{32} \log \frac{1}{32} + \frac{1}{32} \log \frac{1}{32} + \frac{1}{32} \log \frac{1}{32} + \frac{1}{32} \log \frac{1}{32}\right)$$

HUFFMAN CODING

-- OBSERVATIONS (1/2) --

- Observation 1 (about the example):

- **BR=H**, i.e., the coder bitrate achieved the entropy, the best possible
- Does that mean Huffman coding always achieves the entropy?
- No. See the theorem next

- **Theorem:** If all the probabilities (in a memoryless source) are powers of $\frac{1}{2}$, then Huffman achieves $BR=H$. The further away the probabilities are from powers of $\frac{1}{2}$, the further away BR is from entropy H (i.e., $BR > H$).

- We will not prove that theorem, but it is important that you keep it in mind

RUN-LENGTH ENCODING (RLE)

-- RUN-SPLITTING EXAMPLE --

- I: aaaaaaaaaa bbbbbbbbbb dddddddd bbbbbb cccccc aaaaa
- Intermediate code: (a,11) (b,10) (d,8) (b,6) (c,6) (a,5)
- Assume $M=3$ (so, max run-length representable is $7=2^3-1$)
- (a,11) is split into (a,7) (a,4), which in binary is (a,111) (a,100)
- (b,10) is splits into (b,7) (b,3), which in binary is (b,111) (b,011)
- (d,8) is splits into (d,7) (d,1), which in binary is (d,111) (d,001)
- (b,6) (c,6) (a,5) need not be split: in binary (b,110) (c,110) (a,101)
- So, the next intermediate code becomes:
(a,111) (a,100) (b,111) (b,011) (d,111) (d,001) (b,110) (c,110) (a,101)

46

RLE DECODING

- It is an alternating sequence of decodings:
 - Huffman decoding, bin-2-dec of next M bits, Huffman decoding, bin-2-dec of the next M bits, ...
- This decodes to an intermediate representation of (a,L) pairs
- Finally, replace each (a,L) by aaaa...a, where the run is of length L
- Example: Codewords: a:0, b:10, c:110, d:111 $M=3$
 - Bitstream= 01110100101111001111111111001101101101100101
 - Alternating decoding: (a,7) (a,4) (b,7) (b,3) (d,7) (d,1) (b,6) (c,6) (a,5)
 - Final decoding: aaaaaa aaaa bbbbbbb bbb ddddddd d bbbbbb cccccc aaaaa
 - Without spaces: aaaaaaaaaabbbbbbbbbbddddd bbbbbbcccccaaaaa
 - You can verify that the decoded data is identical to the original data

49

GOLOMB CODING

-- PRELIMINARIES (2/2) --

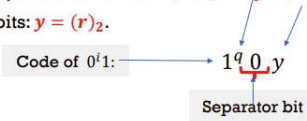
- Golomb coding has a parameter m
 - m is a positive integer set by the algorithm implementer or by the user
 - the optimal value of m = the nearest power of 2 to $p \times \frac{\ln 2}{1-p}$ (proof is later)
 - where p is the probability of the more probable bit in the input stream
 - p is easily computable: count the number of 0's (say N_0) and the number of 1's (say N_1) in the input binary stream, then $p = \max(\frac{N_0}{N_0+N_1}, \frac{N_1}{N_0+N_1})$
 - let's assume that the more probable bit (MPB) is 0
 - If the MPB is 1, then each run is of the form $1^i 0$
- Note: $\ln 2 = 0.6931$

52

GOLOMB CODING

-- METHOD: ASSUME 0 IS THE MPB --

1. Break the input into runs of the form $0^i 1$
2. Code each Golomb run $0^i 1$ as $1^q 0 y$ where
 - Divide i by m , integer division, we get quotient q and a remainder r , i.e., $i = qm + r$
 - y is the binary representation of r , using $\log m$ bits: $y = (r)_2$.



3. The final coded bitstream is: **MPB code₁ code₂ ... code_n tail?** where
 - **MPB**: the (1-bit) value of the more probable bit
 - **code_j**: the code of the j -th Golomb run
 - **tail?**: a single bit that is 1 if the last run has a tail; it is 0 if the last run has no tail

83

GOLOMB CODING

-- EXAMPLE --

- Input stream: $x = 0^9 10^{15} 11$
- MPB=0 since 0 occurs 24 times, while 1 occurs 3 times;
- $p = \frac{24}{27}, p \times \frac{\ln 2}{1-p} = 8 \times 0.6931 = 5.54$, the closest 2-power to 5.54 is 4, so $m = 4, \log m = 2$;
- The Golomb runs of $x = 0^9 10^{15} 11$ are $0^9 1, 0^{15} 1$, and $1 = 0^0 1$
- Code $0^9 1$: $9 = 2 \times 4 + 1$, so $q = 2$ and $r = 1 = (01)_2$, thus code($0^{10} 1$) = $1^2 0 01$
- Code $0^{15} 1$: $15 = 3 \times 4 + 3$, so $q = 3$ and $r = 3 = (11)_2$, thus code($0^{15} 1$) = $1^3 0 11$
- Code $0^0 1$: $0 = 0 \times 4 + 0$, so $q = 0$ and $r = 0 = (00)_2$, thus code($0^0 1$) = $1^0 0 00 = 000$
- tail? = 1 because the last Golomb run has a tail
- The code of x is: $0 1^2 001 1^3 011 000 1$, i.e., 0110011110110001
- $CR = \frac{27}{16} = 1.69$, $BR = \frac{16}{27} = 0.59$ bits/bit

84

Input: a coded bistream (& parameter m); **Output:** the original data

Method:

1. Grab first bit as the MPB, and the last bit of the coded bitstream as the tail
2. Set $k=2$ // index of the bits in the coded bitstream
3. Scan the bitstream rightward from position k looking for successive 1's, keeping a count q , and incrementing k along the way
4. When a 0 is met, read the next $\log m$ bits as y , set $k=k + \log m$
5. Set $r = b2d(y)$ // binary to decimal conversion
6. Compute $i = q \times m + r$
7. Append $0^i 1$ to the output (if MPB=0), else append $1^i 0$ to the output
8. Repeat from 3 until the coded bitstream is exhausted
9. If the last bit (tail?) is 0, strip the final bit from the output

86

GOLOMB DECODING

-- EXAMPLE --

- Example: coded bitstream 0110011110110001, and $m = 4$ ($\log m = 2$)
- MPB=0, tail=1;
- $k=2$, scan the first two 1's till 0, then read 2 bits: 0**11**0011110110001
- So $q=2, y=(01)_2, r=b2d(01)=1, i = q \times m + r = 2 \times 4 + 1 = 9, k=7$
- Append $0^9 1$ to output: output= $0^9 1$
- From $k=7$ scan for 1's till 0: three 1's, 0 and the next two bits 11: 011001**111**0110001
- So $q=3, y=(11)_2, r=b2d(11)=3, i = q \times m + r = 3 \times 4 + 3 = 15, k=13$
- Append $0^{15} 1$ to output: output= $0^9 10^{15} 1$
- Look for 1's from position $k=13$; none found, so $q=0$; skip 0 and read the next 2 bits, $y=00$, so $r=0$; 011001111011**000**1; so $i=0 \times 4 + 0 = 0$; append $0^0 1$: output= $0^9 10^{15} 11$
- Now we reached the last bit, tail=1, so we keep the tail. **Final output: $0^9 10^{15} 11$**

87

DIFFERENTIAL GOLOMB

-- CODING METHOD --

Input: $x = x_1 x_2 \dots x_n$ (& the parameter m)

Output: coded bitstream

Method:

1. Transform x to $z = z_1 z_2 \dots z_n$ where $z_i = x_i - x_{i-1} \forall i > 1$, and $z_1 = x_1$
2. Delete the alternating negatives of the tails, we get $z' = z'_1 z'_2 \dots z'_n$
3. Code z' using Golomb coding

64

DIFFERENTIAL GOLOMB

-- DECODER METHOD --

Input: Coded bitstream (& the parameter m)

Output: Original data

Method:

1. Golomb-decode the coded bitstream into $z' = z'_1 z'_2 \dots z'_n$
2. Keep the first Golomb run's tail at 1
3. Alternate the signs of the remaining tails in z' , getting $z = z_1 z_2 \dots z_n$
4. Set $x_1 = z_1$, and for $i = 2$ to n do: $x_i = x_{i-1} + z_i$
5. Set output = $x_1 x_2 \dots x_n$

65