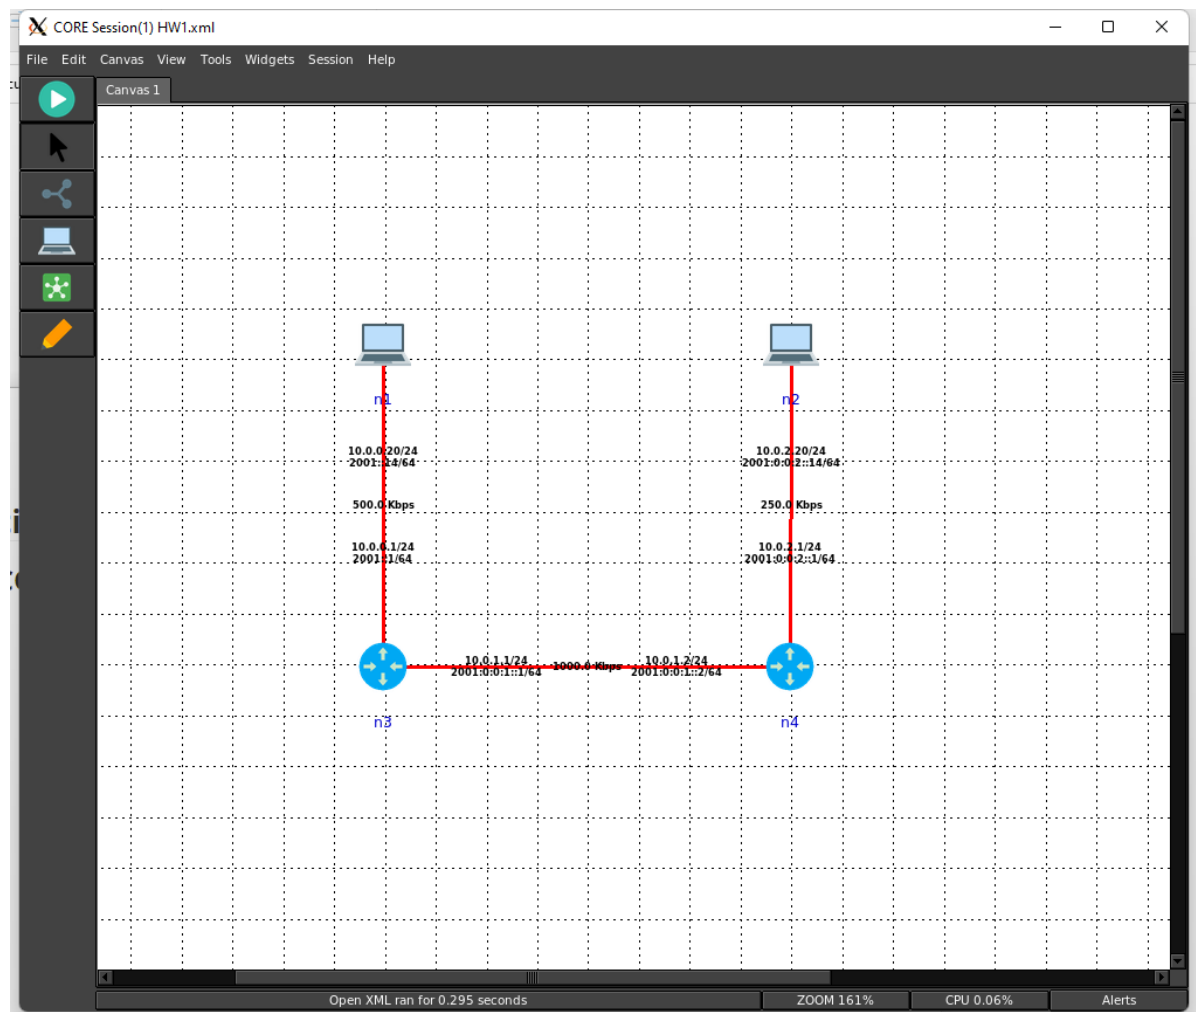


Part 1: Setting up a TCP client and server (2 points)

1) Run the CORE scenario



2) Open a terminal on node n2. Run a server on port 8080:

```
vcmd
root@n2:/tmp/pycore.1/n2.conf# iperf -s -i 1 -p 8080
bind failed: Address already in use
root@n2:/tmp/pycore.1/n2.conf# ps aux|grep iperf
root      34  0.0  0.0   8168   716 pts/2    S+   07:27   0:00 grep --color=
auto iperf
root@n2:/tmp/pycore.1/n2.conf# ps aux|grep nc
root      22  0.0  0.0   3260   820 ?        Ss   07:22   0:00 nc -l 8080
root      36  0.0  0.0   8168   716 pts/2    S+   07:27   0:00 grep --color=
auto nc
root@n2:/tmp/pycore.1/n2.conf# kill 22
root@n2:/tmp/pycore.1/n2.conf# iperf -s -i 1 -p 8080
-----
Server listening on TCP port 8080
TCP window size: 128 KByte (default)
-----
ó
```

- `-s` means run in server mode
- `-i 1` means 1 second between periodic bandwidth reports
- `-p 8080` means server listens on port 8080

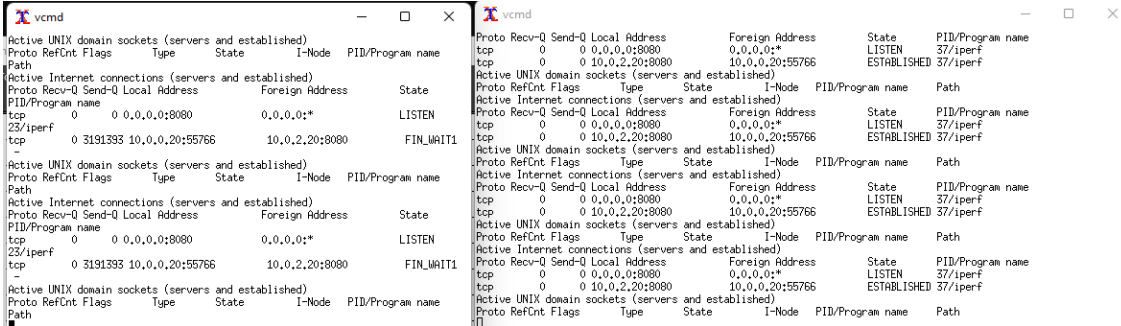
3) Open another terminal on node n2. Use the command `netstat` to list servers

- ```
root@n2:/tmp/pycore.1/n2.conf# netstat -anp
Active Internet connections (servers and established)
Proto Recv-Q Send-Q Local Address Foreign Address State PID/Program name
tcp 0 0 0.0.0.0:8080 0.0.0.0:* LISTEN 37/iperf
Active UNIX domain sockets (servers and established)
Proto RefCnt Flags Type State I-Node PID/Program name Path
root@n2:/tmp/pycore.1/n2.conf#
```
- Proto/Recv-Q/Send-Q/Local Address/Foreign Address/State/PID

### 4) Now run a client on node n1 connecting to node n2

- ```
root@n1:/tmp/pycore.1/n1.conf# iperf -c 10.0.2.20 -i 1 -t 120 -p 8080
-----
Client connecting to 10.0.2.20, TCP port 8080
TCP window size: 85.0 KByte (default)
-----
[ 3] local 10.0.0.20 port 55758 connected with 10.0.2.20 port 8080
[ ID] Interval      Transfer    Bandwidth
[ 3] 0.0- 1.0 sec   107 KBytes  880 Kbits/sec
[ 3] 1.0- 2.0 sec   14.1 KBytes 116 Kbits/sec
[ 3] 2.0- 3.0 sec   36.8 KBytes 301 Kbits/sec
[ 3] 3.0- 4.0 sec   55.1 KBytes 452 Kbits/sec
[ 3] 4.0- 5.0 sec   36.8 KBytes 301 Kbits/sec
[ 3] 5.0- 6.0 sec   42.4 KBytes 348 Kbits/sec
[ 3] 6.0- 7.0 sec   43.8 KBytes 359 Kbits/sec
[ 3] 7.0- 8.0 sec   46.7 KBytes 382 Kbits/sec
[ 3] 8.0- 9.0 sec   42.4 KBytes 348 Kbits/sec
[ 3] 9.0-10.0 sec   46.7 KBytes 382 Kbits/sec
[ 3] 10.0-11.0 sec   46.7 KBytes 382 Kbits/sec
[ 3] 11.0-12.0 sec   45.2 KBytes 371 Kbits/sec
[ 3] 12.0-13.0 sec   43.8 KBytes 359 Kbits/sec
[ 3] 13.0-14.0 sec   42.4 KBytes 348 Kbits/sec
```
- `-c` means run in client mode and connect to 10.0.2.20
 - `-t 120` means set transmit time to 120 seconds


5) Open another terminal on node n1.

- 
- ```

n1:
Active UNIX domain sockets (servers and established)
Proto RefCnt Flags Type State I-Node PID/Program name Path
Active Internet connections (servers and established)
Proto Recv-Q Send-Q Local Address Foreign Address State PID/Program name
tcp 0 0 0.0.0.0:8080 0.0.0.0:* LISTEN 37/iperf
tcp 0 313133 10.0.0.0:8080 0.0.0.0:* LISTEN 37/iperf

n2:
Proto Recv-Q Send-Q Local Address Foreign Address State PID/Program name
tcp 0 0 0.0.0.0:8080 0.0.0.0:* LISTEN 37/iperf
tcp 0 0 10.0.2.20:8080 10.0.0.20:55756 ESTABLISHED 37/iperf

```
- The `recv-Q` and `send-Q` on node 2 is always 0, the `recv-Q` on node1 is always 0, the `send-Q` increase first and drops in the last(takes 240 seconds to clear the `send-Q`, which is exactly 2 times of 120 seconds)



- The Send-Q is the Queue for packet sending, the Recv-Q is the Queue for packet receiving.

## Part 2: Benchmarking (4 points)

### 1) Test 1, throughput:

1. Run the CORE scenario
2. How does iperf measure throughput?
  1. iperf should be counting packet during certain time period to calculate the average throughput
3. Which link in this scenario is the bottleneck link?
  1. The n4->n2 link should be the bottleneck link for
4. Run iperf server on node n2 as shown in part 1
5. Run iperf client on node n1 as shown on part 1
6. The server will report instantaneous throughput per second
7. When the server stops reporting, the last line is the average throughput. What is the average throughput?

```

[4] 234.0-235.0 sec 28.3 KBytes 232 Kbits/sec
[4] 235.0-236.0 sec 29.7 KBytes 243 Kbits/sec
[4] 236.0-237.0 sec 29.7 KBytes 243 Kbits/sec
[4] 237.0-238.0 sec 28.3 KBytes 232 Kbits/sec
1. [4] 238.0-239.0 sec 29.7 KBytes 243 Kbits/sec
 [4] 239.0-240.0 sec 28.3 KBytes 232 Kbits/sec
 [4] 240.0-241.0 sec 29.7 KBytes 243 Kbits/sec
 [4] 0.0-241.5 sec 6.49 MBytes 226 Kbits/sec

```

2. 226 Kbits/sec
8. In your opinion, why is the average throughput smaller than the bottleneck bandwidth?
  1. The bottleneck is just the upper limit of average throughput, there might be packet loss/retry/wait time included to result this smaller average value.
9. Stop the scenario
10. Change rate of link n2-n4 to 750kbps? Repeat starting from (a) to (i)

```


1. 4] 177,0-178,0 sec 58,0 KBytes 475 Kbits/sec
 4] 178,0-179,0 sec 59,4 KBytes 487 Kbits/sec
 4] 179,0-180,0 sec 58,0 KBytes 475 Kbits/sec
 4] 180,0-181,0 sec 58,0 KBytes 475 Kbits/sec
 4] 0,0-181,4 sec 9,43 MBytes 436 Kbits/sec

```

2. The Average increased to 436Kbits/s, Now the bottleNeck Link changed to n1->n3

## 2) Test2,loss

- Reload the scenario (to reset rate of link n2-n4)
- Add a 10% loss on the link n2-n4



- Run the CORE scenario
- Run a ping test from n1- n4 for 1 minute (till you see icmp\_seq=60)
  - Ping 10.0.2.20
  - Control-c to stop
- When you stop ping, it will report statistics. What is the packet loss reported by ping? Can you justify the results (are the results 10% or more? If more, why?)

```

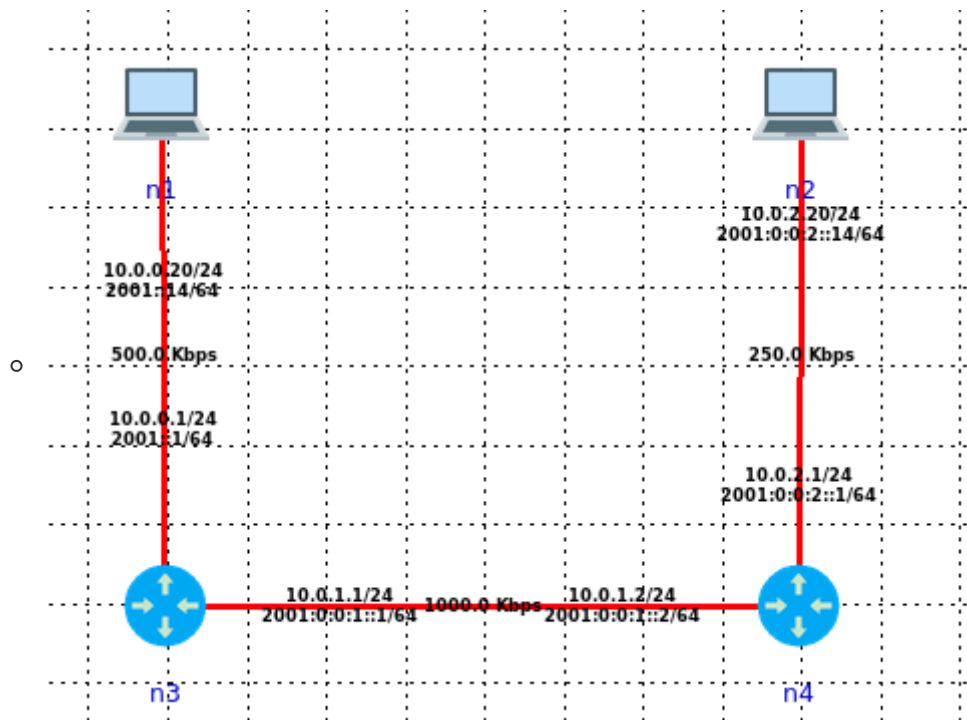
64 bytes from 10.0.2.20: icmp_seq=55 ttl=62 time=11.4 ms
64 bytes from 10.0.2.20: icmp_seq=57 ttl=62 time=11.4 ms
64 bytes from 10.0.2.20: icmp_seq=59 ttl=62 time=11.3 ms
64 bytes from 10.0.2.20: icmp_seq=60 ttl=62 time=11.2 ms
64 bytes from 10.0.2.20: icmp_seq=61 ttl=62 time=11.2 ms
64 bytes from 10.0.2.20: icmp_seq=63 ttl=62 time=11.3 ms
64 bytes from 10.0.2.20: icmp_seq=64 ttl=62 time=11.3 ms
64 bytes from 10.0.2.20: icmp_seq=65 ttl=62 time=11.1 ms
64 bytes from 10.0.2.20: icmp_seq=66 ttl=62 time=11.3 ms
^C
--- 10.0.2.20 ping statistics ---
66 packets transmitted, 51 received, 22.7273% packet loss, time 65377ms
rtt min/avg/max/mdev = 11.143/11.321/11.437/0.065 ms

```

- The loss is about 22.72% and is a lot higher than 10%, the main reason behind this is the loss is applied both to inbound traffic and outbound traffic, resulting a double loss rate in round trip.
- Stop the scenario

### 3) Test3, end to end delay

- a. Read on RTT: [https://developer.mozilla.org/en-US/docs/Glossary/Round\\_Trip\\_Time\\_\(RTT\)](https://developer.mozilla.org/en-US/docs/Glossary/Round_Trip_Time_(RTT)).
- b. Reload the scenario (to reset loss on link n2-n4)
- c. Add a 100ms delay on the link n2-n4



- d. Run the CORE scenario
- e. Run a ping test from n1-n4 for 1 minute (till you see icmp\_seq=60)

```

34 bytes from 10.0.2.20: icmp_seq=49 ttl=62 time=211 ms
34 bytes from 10.0.2.20: icmp_seq=50 ttl=62 time=211 ms
34 bytes from 10.0.2.20: icmp_seq=51 ttl=62 time=211 ms
34 bytes from 10.0.2.20: icmp_seq=52 ttl=62 time=211 ms
34 bytes from 10.0.2.20: icmp_seq=53 ttl=62 time=211 ms
34 bytes from 10.0.2.20: icmp_seq=54 ttl=62 time=211 ms
34 bytes from 10.0.2.20: icmp_seq=55 ttl=62 time=211 ms
34 bytes from 10.0.2.20: icmp_seq=56 ttl=62 time=211 ms
34 bytes from 10.0.2.20: icmp_seq=57 ttl=62 time=211 ms
34 bytes from 10.0.2.20: icmp_seq=58 ttl=62 time=211 ms
34 bytes from 10.0.2.20: icmp_seq=59 ttl=62 time=211 ms
34 bytes from 10.0.2.20: icmp_seq=60 ttl=62 time=211 ms
^C
--- 10.0.2.20 ping statistics ---
30 packets transmitted, 60 received, 0% packet loss, time 59083ms
rtt min/avg/max/mdev = 211.250/211.396/211.534/0.044 ms
root@n1:/tmp/pycore.1/n1.conf#

```

- f. When you stop ping, it will report statistics. What is the RTT (round trip time) reported by ping? Can you justify the results (is delay 100ms or more? If more, why?)
  - The ping's RTT is around 211 ms, is a lot more than 100ms. Because the ping reports RTT which consist outbound time and inbound time. So this 100ms latency is applied twice on the result RTT.
- g. Stop the scenario

## 4) Test4, per link latency

- a. Rerun last scenario (keep 100ms delay on the link n2-n4)
- b. Run a traceroute from n1-n4:
  - `traceroute -n 10.0.2.20`  

```
root@n1:/tmp/pycore.1/n1.conf# traceroute -n 10.0.2.20
traceroute to 10.0.2.20 (10.0.2.20), 30 hops max, 60 byte packets
 1 10.0.0.1 2.920 ms 4.516 ms 6.179 ms
 2 10.0.1.2 7.827 ms 9.475 ms 11.096 ms
 3 10.0.2.20 217.140 ms 220.302 ms 223.646 ms
```
- c. Compare the latency results from ping to traceroute. What information does each provide?
  - Ping mainly provide packet loss and overall RTT
  - Traceroute provides RTT for each node in it's path.
  - Describe a situation when you would want to use ping to measure end to end latency
    - When I want to check if there is packet loss in the link or I'm only interested in overall RTT
  - Describe a situation when you would want to use traceroute to measure per link latency
    - When I want to check latency for each link to my destination