Review Test Submission: Finexam

User	Ruojia Zhang
Course	202203_Database Management Systems_CSCI_6441_10
Test	Finexam
Started	11/27/22 3:16 PM LATE
Submitted	11/27/22 4:54 PM LATE
Due Date	11/21/22 11:59 PM
Status	Needs Grading
Attempt Score	0 out of 20 points
Time Elapsed	1 hour, 37 minutes out of 3 hours
Results Displayed	All Answers, Submitted Answers, Correct Answers, Feedback, Incorrectly Answered Questions

Question 1

0.5 out of 0.5 points

You can select partial table contents by naming the desired fields and by placing restrictions on the rows to be included in the output.

Selected Answer: True

Answers: True

False

Question 2

0.5 out of 0.5 points

Which ALTER TABLE statement should you use to add a PRIMARY KEY constraint on the MANUFACTURER_ID column of the INVENTORY table?

Selected Answer:
ALTER TABLE inventory
ADD PRIMARY KEY (manufacturer_id);

A.

Answers:
ALTER TABLE inventory
ADD PRIMARY KEY (manufacturer_id);

A.

ALTER TABLE inventory
B. ADD CONSTRAINT manufacturer_id PRIMARY KEY;

ALTER TABLE inventory
MODIFY manufacturer_id CONSTRAINT PRIMARY KEY;

C.

ALTER TABLE inventory
MODIFY CONSTRAINT PRIMARY KEY manufacturer_id;

D.

Response Feedback: A - You should use the following ALTER TABLE statement to add a PRIMARY KEY constraint on the MANUFACTURER_ID column of the INVENTORY table:

ALTER TABLE inventory
ADD PRIMARY KEY (manufacturer_id);

To add a PRIMARY KEY constraint to an existing column in a table, use the ALTER TABLE statement with the ADD clause. In this scenario, the ALTER TABLE statement adds a PRIMARY KEY constraint to the MANUFACTURER_ID column in the INVENTORY table. Because a name is not provided for the PRIMARY KEY constraint, a system-generated name will be used. To specify a constraint name, you must use the CONSTRAINT keyword in the ADD clause. For this statement to execute, the table must be empty or all the values in the MANUFACTURER_ID column must be unique and non-null.

The ALTER TABLE statements using the MODIFY clause are incorrect because you cannot add a PRIMARY KEY constraint using this clause. These statements will return a syntax error.

The ALTER TABLE statement containing the clause ADD CONSTRAINT manufacturer_id PRIMARY KEY will return a syntax error because the column name for the PRIMARY KEY constraint is not specified.

Question 3 0 out of 0.5 points

Which statement is NOT CORRECT?

Selected Answer: Only structured data can be described in a formal logical data model.

Answers: To cater for data independence, interfaces between the three layers of the database architecture should be provided.

Physical data independence means that changes in the data storage specifications should not impact the applications, the defined views nor the logical data model.

Logical data independence means that changes in the logical or conceptual data model should not impact the data storage specifications.

Only structured data can be described in a formal logical data model.

Question 4 0.5 out of 0.5 points

You are performing an export of certain objects in your database by using Oracle Data Pump. You are required to unload only the data in the export file and not the metadata information.

Which parameter of the Data Pump Export should you use in this scenario?

Selected Answer: b. CONTENT

Answers:

a. TABLES

b. CONTENT

c. EXCLUDE

d. ESTIMATE

e. DIRECTORY

Response Feedback: You should use the CONTENT parameter of the Data Pump Export. The CONTENT parameter enables you to filter the data, the metadata, or both the data and metadata. The CONTENT parameter has three values: ALL , DATA_ONLY , and METADATA_ONLY . These values specify whether both data and metadata will be exported, only the data will be exported, or only the metadata will be exported.

The TABLES parameter of the Data Pump Export does not allow you to filter the data or metadata that you want to export. The TABLES parameter specifies that you perform a table mode export.

The EXCLUDE parameter of the Data Pump Export does not allow you to filter the data based on the data or metadata that will be exported. The EXCLUDE parameter enables you to exclude specific objects and object types from the metadata that will be exported.

The ESTIMATE parameter of the Data Pump Export determines the mode that export will use to estimate how much space each table in the export job will consume. The estimate done in this process is for the table row data only, not the metadata. This parameter of Data Pump Export is not used to filter the type of data to be exported.

The DIRECTORY parameter of the Data Pump Export specifies the location to which the export writes the dump file set and the log file.

Question 5 0.5 out of 0.5 points

Click the Exhibit(s) button to examine the structure of the CURR_ORDER table. Which aggregate functions could be used with the ORDER_DT column? (Choose all that apply.)

CURR_ORDER

ORDER_ID	NUMBER	NOT NULL, Primary Key
CUSTOMER_ID	NUMBER	Foreign key to CUSTOMER_ID column of the CUSTOMER table
EMP_ID	NUMBER	Foreign key to EMP_ID column of the EMPLOYEE table
ORDER_DT	DATE	
ORDER_AMT	NUMBER(7,2)	
SHIP_METHOD	VARCHAR2(5)	
SHIP_DT	DATE	
STATUS	VARCHAR2(5)	

Selected Answers: MAX
A. MIN
D. COUNT
E.

Answers: MAX
A. SUM
B. AVG
C. MIN
D. COUNT
E. VARIANCE
F.

Response Feedback: A, D, E - The ORDER_DT column in the CURR_ORDER table has a DATE data type. The MAX, MIN, and COUNT group functions can be used with numeric or date data types.

All of the other options are incorrect. Not all group functions can be used with date data types. The SUM, AVG, VARIANCE, and STDDEV group functions can only be used with numeric data types.

Question 6 0.5 out of 0.5 points

True or False? CREATE TABLE is one of the most important SQL Data Manipulation Language (DML) commands.

Selected Answer: False

Answers: True
False

Response Feedback: **Ahead:** Defining the Database: SQL Data Definition Language (DDL)
Complexity: Moderate
Subject: Chapter 5
Title: Relational Database Management Systems and SQL

Question 7 0.5 out of 0.5 points

The following four tables are created and then a view is also created:

SQL> CREATE TABLE georgia
The statement creates a composite non-unique index.

Item: 119 (Ref:1Z0-051.11.1.3)

(id name location credit_limit
number(5), varchar2(25),
varchar2(20), number(9,2));
SQL> CREATE TABLE florida AS SELECT * FROM georgia; SQL> CREATE TABLE alabama AS SELECT * FROM georgia; SQL> CREATE TABLE MISSISSIPPI AS SELECT * FROM georgia;

SQL> CREATE VIEW southeast AS SELECT * FROM georgia

```
UNION
SELECT * FROM florida UNION
SELECT * FROM alabama UNION
SELECT * FROM mississippi;
```

Is it possible to perform DML operations on the VIEW called southeast?

Selected Answer: No, because a view created with the UNION set operator will always be non-updateable.

Answers:
- Yes, since you did not add the clause WITH CHECK OPTION to the CREATE VIEW command.
- Yes, since you did not add the clause WITH READ ONLY to the CREATE VIEW command.
- Yes, because the structures of all four tables are identical.
- No, because a view created with the UNION set operator will always be non-updateable.

Response Feedback: You cannot perform DML operations on the VIEW called southeast because a view created with any set operator (UNION, MINUS, INTERSECT, etc.) will always be a non-updateable view, and hence no DML operations will be possible.

The WITH CHECK OPTION, when added to the CREATE VIEW statement, prohibits the user from inserting or updating a row in the underlying table(s) if that row would be invisible when queried from the view. However, in this command, this is moot since a view created with set operators is never updateable.

The WITH READ ONLY, when added to the CREATE VIEW statement, prohibits the user from inserting, updating, or deleting a row in the underlying table(s). However, in this command, this is moot since a view created with set operators is never updateable.

For tables in a SELECT statement to be combined together by set operators, the tables have the same number of columns and that the columns match in terms of datatype. However, in this command, this is moot since a view created with set operators is never updateable.

Question 8

0.5 out of 0.5 points

The highest normal form for which it is always possible to find a dependency preserving lossless decomposition is:

Selected Answer: D. 3NF.

Answers:
- A. 2NF.
- B. BCNF.
- C. 1NF.
- D. 3NF.

Response Feedback: **Ahead:** Properties of Relational Decompositions
Complexity: Moderate
Subject: Chapter 6
Title: Normalization and Denormalization

Question 9

0.5 out of 0.5 points

All users in the hr_emp role have UPDATE privileges on the employee table. You create the update_employee procedure. The hr_emp users should only be able to update the employee table using this procedure. Which two statements should you execute? (Choose two.)

Selected Answers:
- REVOKE UPDATE ON employee FROM hr_emp;
- C. GRANT EXECUTE ON update_employee TO hr_emp;
- E. GRANT UPDATE ON employee TO hr_emp;

Answers:
- A. GRANT SELECT ON employee to hr_emp;
- B. REVOKE UPDATE ON employee FROM hr_emp;
- C. REVOKE UPDATE ON employee FROM public;
- D.

GRANT EXECUTE ON update_employee TO hr_emp;

E.

Response Feedback: C, E - The two statements you should execute are:

REVOKE UPDATE ON employee FROM hr_emp;
GRANT EXECUTE ON update_employee TO hr_emp;

Unless you are the owner of the PL/SQL construct, you must be granted the EXECUTE object privilege to run it or have the EXECUTE ANY PROCEDURE system privilege. By default, a PL/SQL procedure executes under the security domain of its owner. This means that a user can invoke the procedure without privileges on the procedures underlying objects. To allow hr_emp users to execute the procedure, you must issue the GRANT EXECUTE ON update_employee TO hr_emp; statement.

To prevent hr_emp users from updating the employee table unless they are using the update_employee procedure, you must issue the statement REVOKE UPDATE ON employee FROM hr_emp;

All of the other options are incorrect because they will not meet the specified requirements.

Question 10 0.5 out of 0.5 points

The database contents are loaded during the _____ phase of the Systems Development Life Cycle (SDLC).

Selected Answer: a. implementation

Answers:
a. implementation
b. maintenance
c. analysis
d. detailed systems design

Question 11 0 out of 0.5 points

Which two statements regarding the valid use of single-row and multiple-row subqueries are true? (Choose two.)

Selected Answers: Single-row operators can only be used with single-row subqueries.

C. Multiple-row subqueries can be used in a WHERE clause and the INTO portion of an INSERT statement.

F. Single-row subqueries can only be used in a WHERE clause.

Answers:
A. Multiple-row subqueries can be used with the LIKE operator.
B. Single-row operators can only be used with single-row subqueries.
C. Single- and multiple-row subqueries can be used with the BETWEEN operator.
D. Multiple-row subqueries can be used with both single-row and multiple-row operators.
E. Multiple-row subqueries can be used in a WHERE clause and the INTO portion of an INSERT statement.
F.

Response Feedback: A, F - The following two statements regarding the valid use of single-row and multiple-row subqueries are true:

- Single-row operators can only be used with single-row subqueries.
- Multiple-row subqueries can be used in a WHERE clause and the INTO portion of an INSERT statement.

A single-row subquery is a subquery that returns only one row from the inner SELECT statement. Single-row subqueries can only be used with single-row operators, such as =, >, >=, <, <=, or <>. When a single-row operator is used, then the subquery must be a single-row subquery that returns only one row. If you attempt to use a single-row operator with a subquery that returns multiple rows, an error occurs. Multiple-row subqueries can be used in a WHERE clause and the INTO portion of an INSERT statement.

When used in a WHERE clause, the multiple-row subquery must use a multiple-row operator, such as IN, ANY, or ALL. When used in the INTO portion of an INSERT statement, all rows returned by the multiple-row subquery are inserted into the specified table.

The option stating that single-row subqueries can only be used in a WHERE clause is incorrect. Single-row subqueries can be used in any place that a single scalar value can be used.

The option stating that multiple-row subqueries can be used with the LIKE operator is incorrect. The LIKE operator accepts a single value and can only be used with single-row queries.

The option stating that single- and multiple-row subqueries can be used with the BETWEEN operator is incorrect. The BETWEEN operator accepts two values. Either one or both of these values may come from single-row subqueries. However, the values may not come from multiple-row subqueries.

The option stating that multiple-row subqueries can be used with both single-row and multiple-row operators is incorrect because they can only use multiple-row operators.

Question 12 0.5 out of 0.5 points

For a relational view to be updatable, which of the following must be true?

Selected Answer: C. The view must contain the primary key of the base table it is defined on.

Answers:
A. The view must have all the same attributes as the base table it is defined on.
B. The view table itself must be permanently stored.
C. The view must contain the primary key of the base table it is defined on.
D. All of these must be true.

Response Feedback: **Ahead:** Views
Complexity: Moderate
Subject: Chapter 4
Title: The Relational Model

Question 13 0.5 out of 0.5 points

_____ is the central activity during the parsing phase in query processing.

Selected Answer: d. Query optimization

Answers:
a. Partitioning
b. Query validation
c. Clustering
d. Query optimization

Question 14 0.5 out of 0.5 points

Evaluate this procedure:

```
CREATE OR REPLACE PROCEDURE remove_department
(v_deptno IN NUMBER(9))
IS
BEGIN
DELETE FROM dept
WHERE deptno = v_deptno;
END;
```

Why does this statement fail when compiled?

Selected Answer: B. Specifying a precision for a formal parameter is not permitted.

Answers:
A. The keyword AS is missing.
B. Specifying a precision for a formal parameter is not permitted.
C. V_DEPTNO should be declared as IN OUT.
D. The DELETE statement has a syntax error.
E. The procedure ends without doing a COMMIT or ROLLBACK.

Response Feedback: B - This statement fails because specifying a precision for a formal parameter is not permitted. When declaring a formal argument, only the parameter name and data type must be specified. The data type cannot include precision. You can also specify the mode of the parameter as either IN, OUT, or IN OUT, but this is not required. If the mode is not specified, an IN parameter is assumed.

The option stating that the keyword AS is missing is incorrect because the IS keyword and the AS keyword are interchangeable, and the IS keyword is used in this scenario.

The option stating the argument should be declared as IN OUT is incorrect because the procedure does not pass a value back to the calling environment. IN is the default type of parameter and passes a constant value from the calling environment to the procedure.

The option stating the DELETE statement has a syntax error is incorrect because the statement uses the correct DELETE statement syntax. The option stating that the procedure fails because it doesn't COMMIT or ROLLBACK the transaction is false. The procedure doesn't have to do either, and the status of the current transaction at the conclusion of this procedure is "pending".

Question 15 0.5 out of 0.5 points

In Oracle, the user can force the system to use a particular index by including in the SQL statement:

Selected Answer: B. a + INDEX hint.

Answers:
A. an ORDER BY.
B. a + INDEX hint.
C. an EXPLAIN plan.
D. a join query.

Response Feedback: **Ahead:** Query Optimization in Oracle
Complexity: Easy
Subject: Chapter 10
Title: Relational Query Optimization

Question 16 0 out of 0.5 points

How does the USER_TABLES dictionary view differ from the ALL_TABLES data dictionary view?

Selected Answer: C. ALL_TABLES will display only the tables on which the user has SELECT privileges.

Answers:
A. ALL_TABLES will display only the tables owned by the user.
B. USER_TABLES will display only the tables owned by the user.
C. ALL_TABLES will display only the tables on which the user has SELECT privileges.
D. USER_TABLES will display all the tables on which the user has SELECT privileges.

Response Feedback: B - The USER_TABLES view displays all of the tables that the user owns, while the ALL_TABLES view displays all of the tables to which the user has access.

Question 17 0.5 out of 0.5 points

Code example 4-2

```
SELECT vendor_name, invoice_number
FROM invoices LEFT JOIN vendors
  ON invoices.vendor_id = vendors.vendor_id
```

(Refer to code example 4-2.) If the LEFT keyword is replaced with the RIGHT keyword, the total number of rows that are returned must equal

Selected Answer: d. none of the above

Answers:
a. the number of rows in the Invoices table
b. the number of rows in the Vendors table
c. the number of rows in the Invoices table plus the number of rows in the Vendors table
d. none of the above

Question 18 0.5 out of 0.5 points

The command **SELECT POWER**(4,3) **FROM** DUAL; returns which of the following values?

Selected Answer: D. 64

Answers:
A. 81
B. 4

C. 16

D. 64

Response Feedback: **Ahead:** Additional SQL Functions
Complexity: Moderate
Subject: Chapter 7
Title: Advanced SQL

Question 19 0.5 out of 0.5 points

The _____ constraint assigns a value to an attribute when a new row is added to a table.

Selected Answer: d. DEFAULT

Answers:
a. NOT NULL
b. UNIQUE
c. CASCADE
d. DEFAULT

Question 20 0 out of 0.5 points

You issued the following UPDATE statement, but have not committed the changes:

UPDATE emp SET status='ACTIVE'
WHERE emp_id > 549;
Which statement about the transaction and the corresponding undo data is TRUE?

Selected Answer: D. When your transaction began, it was assigned to multiple undo segments based on the rows being updated.

Answers:
A. Other transactions issued after your UPDATE statement will always be associated with a different undo segment.
B. The statement you issued does not generate any undo data.
C. When your transaction began, it was assigned to an undo segment.
D. When your transaction began, it was assigned to multiple undo segments based on the rows being updated.

Response Feedback: C - When your transaction began, it was assigned to an undo segment. As data is modified, a copy of the "before image" of the updated columns is written to the extents in the assigned segment. When all space in the extent is used, data is written to the next extent in the segment. Each transaction is associated with only one undo segment. Each transaction that is executed while your transaction is running is also assigned to an undo segment, and two transactions might possibly be assigned to the same undo segment.

The option that states other transactions issued after your UPDATE statement will always be associated with a different undo segment is incorrect. Multiple transactions can be assigned to the same undo segment.

The option that states when your transaction began, it was assigned to multiple undo segments based on the rows being updated, is incorrect. A transaction is assigned to a single undo segment.

The option that states the statement you issued does not generate any undo data is incorrect. Undo data is generated when data in the database is modified. This undo data contains before the image of the columns of data being updated. The before image of the data supports read consistency for each user, allowing you to roll back uncommitted transactions, and is used during recovery. Undo data is also required to perform some flashback operations, such as querying the database as it existed at a previous time using Flashback Query or flashing back a transaction.

References:
1. Oracle Database Concepts
12c Release 1 (12.1)
E41396-09
Overview of Segments
https://docs.oracle.com/database/121/CNCPT/logical.htm#CNCPT304
2. Oracle Database Concepts
12c Release 1 (12.1)
E41396-14 Undo Segments
https://docs.oracle.com/database/121/CNCPT/logical.htm#CNCPT305

Question 21 0.5 out of 0.5 points

True or False? When performing a join, the choice of method depends only on the size of the files.

Selected Answer: False

Answers: True
False

Response Feedback: **Ahead:** Physical Query Optimization
Complexity: Difficult
Subject: Chapter 10
Title: Relational Query Optimization

Question 22 0.5 out of 0.5 points

True or False? Object-oriented database management systems allow the database designer to create highly complex, interrelated objects and to give them persistence.

Selected Answer: True

Answers: True
False

Response Feedback: **Ahead:** Rationale for the Object-Oriented (OO) Data Model
Complexity: Easy
Subject: Chapter 9
Title: Object-Based Models

Question 23 0 out of 0.5 points

Which statements concerning the creation of a view are TRUE? (Choose all that apply.)

Selected Answers: View columns that are the result of derived values must be given a column alias.

A view may have column names that are different from the actual base table(s) column names by using column aliases.

Whentheviewalreadyexists,usingtheOR REPLACEoptionrequiresregrantingtheobjectprivilegespreviouslygranted on the view.

Answers: View columns that are the result of derived values must be given a column alias.

A constraint name must be provided when using the WITH CHECK OPTION clause or the statement will fail.

A view may have column names that are different from the actual base table(s) column names by using column aliases.

Whentheviewalreadyexists,usingtheOR REPLACEoptionrequiresregrantingtheobjectprivilegespreviouslygranted on the view.

Question 24 0 out of 0.5 points

When moving data from permanent storage to RAM, an I/O disk operation retrieves a(n):

Selected Answer: a. entire table.

Answers: a. entire table.
b. entire physical disk block.
c. single row containing the attribute requested.
d. single attribute requested only.

Question 25 0.5 out of 0.5 points

In the relational model, relations are represented physically as:

Selected Answer: A. tables.

Answers: A. tables.
B. objects.
C. rectangles.

D. records.

Response Feedback: **Ahead:** Overview of Data Models
Complexity: Easy
Subject: Chapter 2
Title: Database Planning and Database Architecture

Question 26

0.5 out of 0.5 points

You have just moved a stand-alone function into a package. What could you add to the package specification to check the purity level of this function?

Selected Answer: PRAGMA RESTRICT_REFERENCES

Answers:
D. PRAGMA PURITY_CHECK
A. PRAGMA EXCEPTION_INIT
B. PRAGMA FUNCTION_PURITY
C. PRAGMA RESTRICT_REFERENCES
D.

Response Feedback: D - You could add PRAGMA RESTRICT_REFERENCES to the package specification to check the purity level of this function. In prior releases of Oracle, packaged functions must include a guarantee in the package specification that the body of the function does not update the database. PRAGMA RESTRICT_REFERENCES is a directive to the compiler indicating that the package body should not compile correctly if it fails one of its references. Oracle9i and later can read the body of the packaged function during SQL statement processing and does not require this directive.

PRAGMA EXCEPTION_INIT is used to assign exception names.

PRAGMA PURITY_CHECK and PRAGMA FUNCTION_PURITY are not valid compiler directives.

Question 27

0.5 out of 0.5 points

Evaluate this SQL statement:

SELECT line_item_id, order_id, product_id FROM line_item

Which WHERE clause should you include to limit the query results to rows that have a QUANTITY column that has a null value?

Selected Answer: WHERE quantity IS NULL;

Answers:
WHERE quantity = NULL;
WHERE quantity <> NULL;
WHERE quantity != NULL;
WHERE quantity IS NULL;

Response Feedback: You should include the following WHERE clause to limit the query results: WHERE quantity IS NULL;

The IS NULL operator is used to test for the absence of data, also known as null values. When using the IS NULL operator, a Boolean value of TRUE is returned when there is a null value. FALSE is returned when a value exists. The rows returned from this query will include line items that have no QUANTITY value.

All of the other options are incorrect because a NULL value cannot be compared to any value. Comparison operators such as equal (=), greater than (>), less than (<), and not equal to (!= or <>) will not produce the desired result. The value returned usingtheseoperatorstocomparetheQUANTITYcolumntoaNULLvalueisUNKNOWN.TheIS NULLcomparisonoperator tests for NULL values.

Question 28

0.5 out of 0.5 points

All of the following signal an end of an Oracle transaction *except*:

Selected Answer: B. a SAVEPOINT.

Answers:
A. a ROLLBACK statement.
B. a SAVEPOINT.
C. a DDL command.
D. a COMMIT statement.

Response Feedback: **Ahead:** Transaction Management in Oracle
Complexity: Moderate
Subject: Chapter 11
Title: Transaction Management

Question 29 0.5 out of 0.5 points

Which of the following *can* be used with an updatable view?

Selected Answer: a. a join

Answers:
a. a join
b. the UNION operator
c. an aggregate function
d. a DISTINCT clause

Question 30 0.5 out of 0.5 points

You need to add a FOREIGN KEY constraint to the LINE_ITEM_ID column in the CURR_ORDER table to refer to the ID column in the LINE_ITEM table. Which statement should you use to achieve this result?

Selected Answer:
A.
```
ALTER TABLE curr_order
ADD CONSTRAINT line_itemid_fk FOREIGN KEY (line_item_id) REFERENCES line_item (id);
```

Answers:
A.
```
ALTER TABLE curr_order
ADD CONSTRAINT line_itemid_fk FOREIGN KEY (line_item_id) REFERENCES line_item (id);
```
B.
```
ALTER TABLE curr_order
MODIFY (CONSTRAINT lineitem_id_fk FOREIGN KEY (line_item_id)
REFERENCES line_item (id));
```
C.
```
ALTER TABLE line_item
ADD CONSTRAINT curr_order_line_item_id_fk FOREIGN KEY (id)
REFERENCES curr_order (line_item_id);
```
D.
```
ALTER TABLE line_item
MODIFY (CONSTRAINT curr_order_lineitem_fk FOREIGN KEY (id)
REFERENCES curr_order (line_item_id));
```

Response Feedback: A - You should use the following statement:

```
ALTER TABLE curr_order
ADD CONSTRAINT line_itemid_fk FOREIGN KEY (line_item_id) REFERENCES line_item (id);
```

To add a FOREIGN KEY constraint to an existing column in a table, use the ALTER TABLE statement with the ADD clause. In this scenario, the ALTER TABLE command adds a FOREIGN KEY constraint to the LINE_ITEM_ID column in the CURR_ORDER table.This FOREIGN KEY constraint references the ID column in the LINE_ITEM table.

The ALTER TABLE line_item statement is incorrect because it will add a FOREIGN KEY on the ID column in the LINE_ITEM table. The desired result is to add the constraint to the CURR_ORDER table.

The ALTER TABLE statements using the MODIFY clause are incorrect because you cannot add a FOREIGN KEY constraint using this clause. These statements will return a syntax error.

Question 31 0.5 out of 0.5 points

In the relational model, the structure of a relation, including attribute names, domains, constraints, and similar information is called the relation:

Selected Answer: B. schema.

Answers:
A. extension.
B. schema.
C. instance.

D. table.

Response Feedback: **Ahead:** Relational Data Structures
Complexity: Easy
Subject: Chapter 4
Title: The Relational Model

Question 32 0.5 out of 0.5 points

Which logical-level flashback feature gives you the ability to rewind a transaction as if it had never occurred?

Selected Answer: Oracle Flashback Transaction Query

Answers:
C. Oracle Flashback Table
A. Oracle Flashback Version Query
B. Oracle Flashback Transaction Query
C. Oracle Flashback Query
D.

Response Feedback: C - The Oracle Flashback Transaction Query gives you the ability to view transactions and to see their dependencies. Once the dependencies are l then back out these transactions as if the DML did not occur. This process uses Oracle undo data to back out transactions until the database retu state before the DML transaction happened. This functionality is achieved by using the DBMS_FLASHBACK package.

The Oracle Flashback Table feature allows you to recover a table or a set of tables while the database is online. This feature does not allow you to transactions within the database.

The Oracle Flashback Version Query feature gives you the ability to see versions of data based on their metadata. This feature focuses on the dat rather than the transactions associated with the data.

The Oracle Flashback Query feature allows you to view a database at a specific time in the past. For example, this would give you the ability to rec DML operation occurred and you needed to back it out.

References:

Database Backup and Recovery User's Guide
12c Release 2 (12.2)
E85627-03
Logical Flashback Features
https://docs.oracle.com/en/database/oracle/oracle-database/12.2/bradv/introduction-backup-recovery.html#GUID-725A5589-5ED5-4C02-A0B9-D

Question 33 0.5 out of 0.5 points

The basic objects in the ER model include all of the following *except*:

Selected Answer: D. classes.

Answers:
A. relationships.
B. entities.
C. attributes.
D. classes.

Response Feedback: **Ahead:** Purpose of the Entity-Relationship (ER) Model
Complexity: Moderate
Subject: Chapter 3
Title: The Entity-Relationship Model

Question 34 0.5 out of 0.5 points

A design trap occurs when a relationship is improperly or incompletely identified and is therefore represented in a way that is not consistent with the real world.

Selected Answer: True

Answers: True
False

Question 35 0.5 out of 0.5 points

The Crow's Foot notation easily identifies multivalued attributes.

Selected Answer: False

Answers: True
False

Question 36 0.5 out of 0.5 points

Which of the following isn't true when you're using SQL Developer to test a script that contains more than one SQL statement?

Selected Answer: a. When you run a script, the results are displayed in the Results tab.

Answers: a. When you run a script, the results are displayed in the Results tab.

b. You can run each statement in the script by itself.

c. You can run all of the statements in the script at once.

Question 37 0.5 out of 0.5 points

True or False? Using JDBC, a CallableStatement object is used to execute a stored procedure.

Selected Answer: True

Answers: True
False

Response Feedback: **Ahead:** Java Database Connectivity (JDBC)
Complexity: Moderate
Subject: Chapter 7
Title: Advanced SQL

Question 38 0.5 out of 0.5 points

How does a client read a file from HDFS?

Selected Answer: The client sends a request to the NameNode. The NameNode will return the blocklocations of which DataNode(s) contain the desired information. The client then reads the data directly of the DataNode(s).

Answers: The client queries all DataNodes in parallel. The DataNode that contains the requested data responds directly to the client, while the client reads the data directly off the DataNode.

The client sends a request to the NameNode. The NameNode will return the blocklocations of which DataNode(s) contain the desired information. The client then reads the data directly of the DataNode(s).

The client sends a request to the NameNode. The NameNode then queries the DataNodes for the blocklocations. The DataNodes respond to the NameNode and the NameNode will pass on the information to the client. The client then reads the data directly of the DataNode(s).

The client sends a request to the NameNode. The NameNode contacts the DataNode(s) that contain the desired information. The DataNode(s) send the data to the NameNode, who forwards it the client.

Question 39 0.5 out of 0.5 points

You need to add a FOREIGN KEY constraint to the LINE_ITEM_ID column in the CURR_ORDER table to refer to the ID column in the LINE_ITEM table.

Which statement should you use to achieve this result?

Selected Answer:
```
ALTER TABLE curr_order
ADD CONSTRAINT line_itemid_fk FOREIGN KEY (line_item_id) REFERENCES line_item (id);
```

Answers:
```
ALTER TABLE curr_order
ADD CONSTRAINT line_itemid_fk FOREIGN KEY (line_item_id) REFERENCES line_item (id);
```

```
ALTER TABLE curr_order
MODIFY (CONSTRAINT lineitem_id_fk FOREIGN KEY (line_item_id) REFERENCES line_item (id));
```

```
ALTER TABLE line_item
ADD CONSTRAINT curr_order_line_item_id_fk FOREIGN KEY (id) REFERENCES curr_order (line_item_id);


ALTER TABLE line_item
MODIFY (CONSTRAINT curr_order_lineitem_fk FOREIGN KEY (id)
REFERENCES curr_order (line_item_id));
```

Response Feedback: You should use the following statement:

```
ALTER TABLE curr_order
ADD CONSTRAINT line_itemid_fk FOREIGN KEY (line_item_id) REFERENCES line_item (id);
```

To add a FOREIGN KEY constraint to an existing column in a table, use the ALTER TABLE statement with the ADD clause. In this scenario, the ALTER TABLE command adds a FOREIGN KEY constraint to the LINE_ITEM_ID column in the CURR_ORDERtable.ThisFOREIGN KEYconstraintreferencestheIDcolumnintheLINE_ITEMtable.

The ALTER TABLE line_item statement is incorrect because it will add a FOREIGN KEY on the ID column in the LINE_ITEM table. The desired result is to add the constraint to the CURR_ORDER table.

TheALTER TABLEstatementsusingtheMODIFYclauseareincorrectbecauseyoucannotaddaFOREIGN KEYconstraint using this clause. These statements will return a syntax error.

Question 40 0.5 out of 0.5 points

If S = { a,b} and T = {1,2,3}, then the Cartesian product S X T has how many ordered pairs?

Selected Answer: C. 6

Answers:
A. 5
B. 9
C. 6
D. 8

Response Feedback: **Ahead:** Relational Data Structures
Complexity: Moderate
Subject: Chapter 4
Title: The Relational Model

Sunday, November 27, 2022 4:54:47 PM EST

← OK